

handoff

Describe it. Hand it off. It runs.

A voice-first agent on the **Strands Agents SDK** that runs your recurring chores on AWS and asks you only when it genuinely cannot decide.

Multi-App AI Agent Hackathon · Sept 13, 2026

8 external apps

Bedrock · AgentCore · Transcribe · Polly

MIT

Built for the Multi-App AI Agent Hackathon

Virtual · September 13, 2026 — hosted by **Lemma** and **Comma Capital**, judged by the founders of **Arga Labs** and **Userlens**.

The brief: build one useful, multi-step AI agent, connect it to at least three external apps, and show how you know it works.

The brief, met

- One agent, six-plus steps per run
- **Eight** external apps, not three — Gmail, Google Calendar, Linear, Slack, GitHub, Notion, Telegram, Airtable
- Reliability shown as a table, not a claim: 313 automated tests, handoff doctor making a real call per credential, every decision logged to Airtable with its confidence

Judged on

- Technical execution — 30%
- Reliability & evaluation — 25%
- Usefulness — 20%
- Originality — 15%
- Demo clarity — 10%

The half hour nobody gets back

Every professional starts the day with the same twenty minutes of small work: triage the inbox, file the real asks, archive the noise, reply to the manager.

- Too small to automate. Too constant to ignore.
- Every tool built to help is *another thing to open*.
- And the one thing you actually fear from an agent: it does something you didn't sanction.


The brief asked for an agent that runs autonomously and only surfaces when there's a real decision to make. We built that sentence.

Say it once


Tap the orb and say:

"Every weekday at eight, triage my inbox. Real asks become Linear tickets, newsletters get archived, ask me about anything unsure. Set it up and run it now."

Four tool calls later it has picked the integrations, written the cron, drawn the **line** between what it may do alone and what it must ask about — and switched it on.



It runs, and mostly leaves you alone



The Strands Graph, drawn live from its own hook events. Seven handled alone. One set aside.

The gate — the whole product in one line

```
response = event.interrupt(interrupt_name, reason=payload.model_dump(mode="json"))
```

- A `BeforeToolCallEvent` hook in front of the **one tool that changes anything**.
- Below the confidence threshold it raises *before the tool body runs* — nothing has happened to your mail.
- Unsure items are **deferred, not interrupted**: finish the pass, then ask **once**.
- The graph state is serialised: a run paused at 08:04 and answered at 11:30 is the same run.


Amber means exactly one thing in Handoff: waiting on you.

One screen, one decision — answered out loud

The agent's reasoning, verbatim: *"Unknown vendor pitching a strategic partnership with a CEO intro call. I can't tell if this is a legitimate vendor or phishing."*

40 % sure — needs 70 % to act alone.

Say *"archive it, it's cold outreach."* Matched locally, no model round-trip. The run resumes where it stopped.



It gets quieter

Your answer becomes a **narrow rule**: exact sender, whole domain, or two distinct keywords.
Never one shared word.

Run	Handled alone	From your rules	Waiting on you
Morning 1	7	0	1
Morning 2	7	1	0

Verified on Bedrock Nova. The learner's restraint matters more than its reach.

Built on Strands, all the way down

SDK feature	Where it earns its place
<code>Graph</code> + <code>GraphBuilder</code>	trigger → executor → completer, deterministic, replayable
<code>BeforeToolCallEvent</code> + <code>event.interrupt()</code>	the gate; resume returns the human's answer
<code>BeforeInvocationEvent</code> / <code>AfterInvocationEvent</code>	a narrator per node — the orb draws the Graph live
<code>SessionRepository</code> + <code>RepositorySessionManager</code>	chats persisted in DynamoDB; terminal and browser share one session
<code>@tool</code> s + <code>MCPCClient</code>	Gmail, Google Calendar, Linear, Slack, GitHub, Notion, Telegram, Airtable, web
Context variables into tools	<code>activate_workflow</code> and <code>start_run</code> reach the page

On AWS, for real

- **Amazon Bedrock** — Nova Pro / Nova Lite through cross-region profiles
- **AgentCore Runtime** — arm64 container, session per run, long-running invocations
- **AgentCore Memory** — learned preferences across runs
- **Amazon Transcribe** streaming + **Amazon Polly** — the voice, over a WebSocket
- **DynamoDB** single table · **EventBridge Scheduler** → **Lambda** bridge · **ECR**

Every store has a local JSON backend too, so a fresh clone runs the whole loop offline.

Three surfaces, one gate

Talk — the orb, hands-free, Ctrl+Space anywhere

Desktop — native window, records the microphone itself

CLI — every feature, `--json` everywhere

```
$ handoff run inbox-triage --watch
$ handoff pending
$ handoff decide int_... archive
$ handoff talk --mic
```

```
→ handoff (main) $ handoff run inbox-triage-morning --watch
18:20:46 fetched      Fetched 8 unread messages (synthetic inbox)
18:20:46 tool_end     fetch_unread_emails ok
18:20:51 tool_start   calling submit_action
18:20:51 acted        Filed a ticket for "Need API review by Thursday" - 90% sure
18:20:51 tool_start   calling submit_action
18:20:51 acted        Archived "TC Daily: AI Agents Are Eating Software" - 80% sure
18:20:51 tool_start   calling submit_action
18:20:51 acted        Drafted a reply to "Quick sync on Q3 targets" - 90% sure
18:20:51 tool_start   calling submit_action
18:20:51 deferred     Not sure about "Strategic partnership opportunity - time sensitive" (60%) - setting it aside fo
18:20:51 tool_start   calling submit_action
18:20:51 acted        Left alone "Updated PTO policy - action may be required" - 70% sure
18:20:51 tool_start   calling submit_action
18:20:51 acted        Filed a ticket for "Re: Sprint retro action items" - 90% sure
18:20:51 tool_start   calling submit_action
18:20:51 acted        Filed a ticket for "[handoff/handoff] Dependabot alert: lodash prototype pollution" - 90% sure
18:20:51 tool_start   calling submit_action
18:20:51 tool_start   calling finish_batch
18:20:51 memory       Archived "FYI: Budget freeze through end of month" (a rule you set) - 80% sure
18:20:51 tool_end     submit_action ok
18:20:51 tool_end     submit_action ok
18:20:51 tool_end     submit_action ok
18:20:51 tool_end     submit_action ok
18:20:51 tool_end     submit_action ok
18:20:51 tool_end     submit_action ok
18:20:51 tool_end     submit_action ok
18:20:51 tool_end     submit_action ok
18:20:51 node_end     executor done
[handoff:console] Handoff needs your input
Item: Strategic partnership opportunity - time sensitive - from Unknown - Vendor Corp
From: partnerships@vendor.com
Why: Confidence (60%) is below the 70% threshold
My best guess: skip (60% confidence)
Decide: http://localhost:8000/decide/int_3990c60b15be
18:20:51 asked        Set aside 1 item for you
run          run_826bda5c0873
status       waiting_on_human
handled alone 6
decided by you 0
from your rules 1
summary      Paused on 1 decision that need you.
! 1 decision(s) need you - 'handoff pending'
```

Who it's for

- **The consultant** whose inbox is the business: asks become tickets, the rest is archived, unknown vendors get asked about.
- **The engineering lead:** PR review triage every morning, blocked tickets surfaced, one Slack memo.
- **The founder:** a weekly competitor pricing watch that says nothing when nothing moved.
- **Anyone with a chore on a schedule** — described once, in a sentence.

Why not just ask a chatbot?

A chatbot waits for you. Handoff **runs when you're not there**, on a schedule, with a written line it will not cross, an audit trail that says who decided what, and a memory that makes it ask less every week.

The hard question for a background agent is *when may it act alone?* Everything here is one answer to that question — and the answer moves as it learns.

handoff

**It does the boring part.
And it knows when to stop and ask.**

Multi-App AI Agent Hackathon · github.com/ZINKUNO/Handoff · handoff-eya.pages.dev · MIT